

# Reinforcement Learning Markov Decision Process Notes

Lu Hong <sup>1</sup>

<sup>1</sup> Nanjing University of Aeronautics and Astronautics

September 22, 2019

## 1 Markov Process

"The future is independent of the past given the present". A state  $S_t$  is *Markov* iif.

$$\mathbb{P}[S_{t+1}|S_t] = \mathbb{P}[S_{t+1}|S_1, S_2, \dots, S_t]$$

This requires environment to be fully observed. Markov Process (or Markov Chain) is a *memoryless random process*, represented by a tuple  $\langle S, P \rangle$

### 1.1 Markov Chain

suppliment from WangCai's Note

Let  $s_t \in S$ , and if  $S$  is countable, we call the Markov Process with this countable state space Markov Chain.

## 2 Markov Reward Process

Markov Reward Process is Markov Process with values, represented by a tuple  $\langle S, P, R, \gamma \rangle$ , where  $R = \mathbb{E}[R_{t+1}|S_t = s]$  and  $\gamma$  is a discount factor.

Return  $G_t$  is the *total discounted reward* at time-step  $t$  presented by

$$G_t = \sum_{i=0}^{\infty} \gamma^i R_{t+i+1}$$

Also, we define Value Function to indicate the long-term value of the state  $s$ .


$$V(s) = \mathbb{E}[G_t|S = s]$$

### 2.1 Bellman Equation

\* Bellman Equation for MRPs, it demonstrate that MRPs can be presented in recursive format.

Bellman Equation for MRPs (2)

$$v(s) = \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) | S_t = s]$$



$$v(s) = R_s + \gamma \sum_{s' \in S} P_{ss'} v(s')$$

$$\begin{aligned} v(s) &= \mathbb{E}[G_t | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \dots) | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) | S_t = s] \end{aligned}$$

To express Bellman Equation in matrix form as

$$v = R + \gamma P v,$$

we can solve Bellman Equation easily as linear equation.

$$\begin{aligned} v &= R + \gamma P v \\ (I - \gamma P)v &= R \\ v &= (I - \gamma P)^{-1} R \end{aligned}$$

•

### 3 Markov Decision Process

A MDP is a MRP with decisions, represented by a tuple  $\langle S, A, P, R, \gamma \rangle$ , to be specific, some params have been changed after action is introduced.

$$P_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$$

$$R_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$

#### 3.1 Policies

Policies: A policy  $\pi$  is a distribution over actions given states,

$$\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$$

- A policy fully characterizes the agent's behaviour
- Policy is what we want in RL problem
- Policy is only associated with current state
- Policy is static if it's certain
- Agent can update policy during the time
- If  $\pi$  is one-hot, then policy is certain

#### Policies (2)

- Given an MDP  $\mathcal{M} = \langle S, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and a policy  $\pi$
- The state sequence  $S_1, S_2, \dots$  is a Markov process  $\langle S, \mathcal{P}^\pi \rangle$
- The state and reward sequence  $S_1, R_2, S_2, \dots$  is a Markov reward process  $\langle S, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$
- where

$$\mathcal{P}_{s,s'}^\pi = \sum_{a \in \mathcal{A}} \pi(a|s) \mathcal{P}_{ss'}^a$$

$$\mathcal{R}_s^\pi = \sum_{a \in \mathcal{A}} \pi(a|s) \mathcal{R}_s^a$$

After we introduce Policy, we should reinvent the value function. the state-value function is:

$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s]$$

the action-value function is:

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$$

#### 3.2 Bellman Expectation Equation

##### Bellman Expectation Equation


The state-value function can again be decomposed into immediate reward plus discounted value of successor state,

$$v_\pi(s) = \mathbb{E}_\pi[R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s]$$

The action-value function can similarly be decomposed,

$$q_\pi(s, a) = \mathbb{E}_\pi[R_{t+1} + \gamma q_\pi(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$


##### Bellman Expectation Equation for $V^\pi$



$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) q_\pi(s, a)$$

For each action nod, q-value is generated the same way.


##### Bellman Expectation Equation for $Q^\pi$



$$q_\pi(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_\pi(s')$$


Two process combined in various order, we can get Bellman Expectation Equation for  $v_\pi$  and  $q_\pi$

##### Bellman Expectation Equation for $v_\pi$ (2)



$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_\pi(s') \right)$$

### Bellman Expectation Equation for $q_\pi$ (2)



$$q_\pi(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a'|s') q_\pi(s', a')$$

### 3.3.3 Find Optimal Policy

#### Optimal Value Functions Finding an Optimal Policy

An optimal policy can be found by maximising over  $q_*(s, a)$ ,

$$\pi_*(a|s) = \begin{cases} 1 & \text{if } a = \operatorname{argmax}_{a \in \mathcal{A}} q_*(s, a) \\ 0 & \text{otherwise} \end{cases}$$

- There is always a deterministic optimal policy for any MDP
- If we know  $q_*(s, a)$ , we immediately have the optimal policy

## 3.3 Optimal Value Function

### 3.3.1 Definition

The optimal state-value function  $v_*(s)$  is the maximum value function over all policies

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$


The optimal action-value function  $q_*(s, a)$  is the maximum action-value function over all policies

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

## 3.4 Bellman Optimality Equation

### Bellman Optimality Equation for $v_*$

The optimal value functions are recursively related by the Bellman optimality equations:



$$v_*(s) = \max_a q_*(s, a)$$

### 3.3.2 Optimal Policy

#### Optimal Policy

Define a partial ordering over policies


$$\pi \geq \pi' \text{ if } v_{\pi}(s) \geq v_{\pi'}(s), \forall s$$

#### Theorem

For any Markov Decision Process


- There exists an optimal policy  $\pi_*$  that is better than or equal to all other policies,  $\pi_* \geq \pi, \forall \pi$
- All optimal policies achieve the optimal value function,  $v_{\pi_*}(s) = v_*(s)$
- All optimal policies achieve the optimal action-value function,  $q_{\pi_*}(s, a) = q_*(s, a)$

### Bellman Optimality Equation for $Q^*$



$$q_*(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s')$$

## Bellman Optimality Equation for $V^*$ (2)



$$v_*(s) = \max_a \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s')$$

, We can extract this process into a matrix:

Reward	-1	-2	-2	-2	10	1	0
State	FaceBook	Class1	Class2	Class3	Pass	Pub	Sleep
FaceBook	0.9	0.1					
Class1	0.5		0.5				
Class2				0.8			0.2
Class3					0.6	0.4	
Pass		0.2	0.4	0.4			1
Pub							
Sleep							1

where

$$\mathcal{S} = \begin{bmatrix} \text{FaceBook} \\ \text{Class1} \\ \text{Class2} \\ \text{Class3} \\ \text{Pass} \\ \text{Pub} \\ \text{Sleep} \end{bmatrix},$$

$$\mathcal{P} = \begin{bmatrix} 0.9 & 0.1 & & & & & \\ 0.5 & & 0.5 & & & & \\ & & & 0.8 & & & 0.2 \\ & & & & 0.6 & 0.4 & \\ & 0.2 & 0.4 & 0.4 & & & 1 \\ & & & & & 1 & \\ & & & & & & 1 \end{bmatrix}$$

$$\mathcal{R} = \begin{bmatrix} -1 \\ -2 \\ -2 \\ -2 \\ 10 \\ 1 \\ 0 \end{bmatrix},$$

$$\gamma \in [0, 1]$$

where we can use these params to calculate state value  $v(s)$  using bellman equation

$$v = (I - \gamma P)^{-1} R$$

$$= \left( \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix} - \gamma \begin{bmatrix} 0.9 & 0.1 & & & & & \\ 0.5 & & 0.5 & & & & \\ & & & 0.8 & & & 0.2 \\ & & & & 0.6 & 0.4 & \\ & 0.2 & 0.4 & 0.4 & & & 1 \\ & & & & & 1 & \\ & & & & & & 1 \end{bmatrix} \right)^{-1} \begin{bmatrix} -1 \\ -2 \\ -2 \\ -2 \\ 10 \\ 1 \\ 0 \end{bmatrix}$$

The result state-value vector is stored in  $s$ , shown in 1

```
>> R = [-1;-2;-2;-2;10;1;0]
>>
>> inv(E + E*gamma)
ans =
-22.5427
-12.5429
1.4568
4.3210
10.0000
0.8025
0
```




Figure 1: MRP results

## 4 Example Demonstration


I import example from (Silver, 2018) and explanation from (WangCai, 2019), accompanied with my personal perspective to finish this part.

### 4.1 MRP Example

Given the following Markov Reward Process


MRP

Example: Student MRP



### 4.2 MDP Example

Introduce MDP example from the lecture.



First, we can use state-value function to calculate the value of each state. Using

$$v_{\pi}(s) = \sum_{a \in A} \pi(a|s) (R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_{\pi}(s'))$$

since  $\pi(a|s) = 0.5$ , it means the decision is uniformly random. Also, the state value is initialized as zero-vector actually. for  $v(1)$ ,

$$v(1) = 0.5 * (-1 + v(4)) + 0.5 * (-2 + v(2))$$

for  $v(2)$ ,

$$v(2) = 0.5 * (-2 + v(3))$$

for  $v(3)$ ,

$$v(3) = 0.5 * (10 + 0) + 0.5 * (1 + (0.2 * v(1) + 0.4 * v(2) + 0.4 * v(3)))$$

for  $v(4)$ ,

$$v(4) = 0.5 * (-1 + v(4)) + 0.5 * (0 + v(1))$$

From these four equations, we can calculate four state values.

$$\begin{bmatrix} v1 \\ v2 \\ v3 \\ v4 \end{bmatrix} = \begin{bmatrix} -1.3 \\ 2.7 \\ 7.4 \\ -2.3 \end{bmatrix}$$

After we have got state value, we can use the same method to obtain the value of action.

## Reference

- Silver, D. (2018). "Reinforcement Learning". In: *Teaching*. URL: <http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>.
- WangCai (2019). In: URL: <https://zhuanlan.zhihu.com/p/50685812>.